

VarBlockSpMM: Supplementary Evidence

Cohen Brunner Thomas

6 September 2026

1 Original NVIDIA controls

This table and figure use the original seed-1, three-process campaign. The comparator set includes cuSPARSE and grouped cuBLAS, without MAGMA. They remain separate from the focused MAGMA follow-up in the main paper.

Table 1: Original width breakdown, with process-averaged ratios.

RHS	CSR	Grouped	BSR8	Best library	Direct wins	Min
8	3.513	9.116	1.078	1.057	26/32	0.951
16	4.415	7.778	1.086	1.042	24/32	0.823
32	4.100	4.451	1.128	1.087	21/32	0.903
64	5.190	3.481	1.444	1.397	32/32	1.112




Figure 1: Original core width comparison. Each point averages three process ratios. The fastest-library panel uses the original NVIDIA controls.

Table 2: Descriptive process-label ranges on the core grid. The two endpoints are observed aggregate ratios, not confidence limits.

Comparator	Geometric mean	Process-label range
CSR default	6.834	6.830 to 6.838
CSR lower envelope	4.262	4.253 to 4.268
Grouped cached	5.757	5.755 to 5.759
Grouped changing	5.752	5.749 to 5.754
Subdivided BSR8	1.175	1.174 to 1.176

Table 3: Width and degree breakdown for the core grid. Times are milliseconds and speedups are dimensionless.

RHS	Degree	Direct ms	CSR ratio	Grouped ratio	BSR8 ratio
8	2	0.114	2.643	7.839	1.105
8	4	0.207	3.667	8.901	1.081
8	8	0.386	5.598	9.741	1.068
8	16	0.751	2.807	10.158	1.058
16	2	0.145	3.553	6.244	1.052
16	4	0.248	4.140	7.577	1.069
16	8	0.447	5.643	8.647	1.109
16	16	0.873	4.576	8.945	1.117
32	2	0.283	3.456	3.461	1.086
32	4	0.508	3.655	4.073	1.039
32	8	0.836	4.476	5.083	1.177
32	16	1.582	4.998	5.480	1.218
64	2	0.415	4.619	2.731	1.445
64	4	0.727	4.995	3.416	1.434
64	8	1.352	5.278	3.825	1.436
64	16	2.588	5.960	4.116	1.461

Table 4: Size and degree-pattern controls with random column locality. Each row averages eight configurations and three process observations.

Block rows	Degree pattern	CSR ratio	Grouped ratio	CSR wins
256	16 in every row	2.870	6.428	8/8
256	0, 1, 16, 16	1.492	4.826	8/8
1024	16 in every row	3.056	5.552	8/8
1024	0, 1, 16, 16	3.677	5.294	8/8
4096	16 in every row	4.082	6.224	8/8
4096	0, 1, 16, 16	4.763	5.629	8/8

Table 5: Matched control time divided by changed-policy time. RHS-8 effects are geometric marginal contrasts over the other two factorial choices. Minima and maxima first average three process ratios for each fixed matrix instance. They are observed ranges, not confidence limits.

	RHS	Change	Ratio	Min	Max
	8	Eight-row mapping	1.075	0.995	1.153
	8	Two accumulators	0.972	0.825	1.084
	8	One CTA instead of two	0.943	0.837	1.011
	8	Full-panel staging	1.472	1.120	1.966
	16	Wider accumulation	1.084	0.918	1.506
	16	Cooperative staging	1.406	1.102	1.897
	16	Double buffering	1.022	0.925	1.104
	16	Full-panel staging	2.271	1.866	2.997
	32	Wider accumulation	1.310	1.113	1.598
	32	Cooperative staging	1.414	1.205	1.770
	32	Double buffering	1.012	0.846	1.191
	64	Wider accumulation	1.379	1.217	1.731
	64	Cooperative staging	1.632	1.493	1.857
	64	Double buffering	0.801	0.661	0.924




Figure 2: GPU time ratios across the 128 core configurations. Each ratio is the geometric mean of three paired process ratios, computed from per-process timing medians. Each curve sorts configurations independently, so equal ranks need not identify the same matrix. The vertical axis is logarithmic. Ratios above the dashed parity line favor direct execution. These are empirical distributions over the chosen grid, not uncertainty intervals.

Table 6: All published-matrix controls, averaged across twelve matrix and width configurations. Ratios divide comparator GPU time by direct time.

Comparator	Ratio	Minimum	Direct wins
Compact CSR	0.352	0.210	0/12
BSR8	0.648	0.498	0/12
Grouped cached	5.246	3.361	12/12
Dense SGEMM	1.051	0.460	7/12

Table 7: Near-field comparator GPU time divided by direct time. Each row averages three geometry seeds and three processes per seed. Ratios above one favor direct execution.

Points	Geometry	RHS	CSR	BSR8	Grouped	Dense
2048	uniform	8	0.466	0.371	19.250	0.136
2048	uniform	16	0.557	0.276	13.892	0.121
2048	uniform	32	1.070	0.300	14.079	0.166
2048	uniform	64	0.871	0.252	9.183	0.124
2048	clustered	8	1.012	0.502	16.778	0.163
2048	clustered	16	1.188	0.382	12.269	0.165
2048	clustered	32	1.581	0.286	8.679	0.150
2048	clustered	64	1.754	0.425	7.417	0.129
8192	uniform	8	0.535	0.372	22.483	1.185
8192	uniform	16	0.620	0.253	13.837	1.412
8192	uniform	32	1.032	0.325	11.958	0.490
8192	uniform	64	1.079	0.304	6.529	0.263
8192	clustered	8	1.181	0.598	13.678	0.816
8192	clustered	16	1.836	0.513	10.906	1.256
8192	clustered	32	2.350	0.668	7.097	0.331
8192	clustered	64	3.190	0.836	4.947	0.224

2 Transport discretization and follow-up controls

For element width h , the modal mass matrix has entries $M_{ii} = h/(2i + 1)$. On the reference interval, define $Q_{ij} = \int_{-1}^1 P'_i(x)P_j(x) dx$, which equals two when $i > j$ and $i - j$ is odd and zero otherwise. The periodic upwind transport operator has blocks

$$L_{ij}^{\text{self}} = \frac{2i + 1}{h}(Q_{ij} - 1), \quad L_{ij}^{\text{upwind}} = \frac{2i + 1}{h}(-1)^i.$$

The stored matrix is $I + \Delta t L$, with $\Delta t = 0.05h/64^2$. Element polynomial orders repeat 7, 15, 31, 63. Their unequal modal counts produce the variable block sizes directly. All entries in each stored block are nonzero for the chosen timestep. Initialization projects $\sin(2\pi kx)$ for $k = 1, \dots, N$ using 96-point Gauss-Legendre quadrature. The analytic comparison projects $\sin(2\pi k(x - 32\Delta t))$ and weights coefficient errors by $h/(2i + 1)$. The reported short-trace check is not a long-time stability or convergence study.

The trace comparator starts from identical prepared host blocks. Startup includes method construction and one complete trace, including the reset of its initial device buffer. Common panel allocations and Python assembly are excluded. Every method's output is checked against repeated FP64 sparse products using the assembled FP32 coefficients. No CPU reference work occurs inside timing.

MAGMA slot execution needs no product workspace. The reduction composition allocates $64N$ FP32 entries per stored block. The alternating-buffer transport trace retains two pointer plans and thus twice that product workspace. The reduction kernel and all repeated slot calls are timed. All MAGMA code is pinned in the source archive. This composition comparison does not reproduce the full hierarchical solver from the motivating work.

Table 8: Controlled combinations. Best-library-to-direct ratios average three paired process ratios. Fixed-size 32 cases include BSR32.

Rows	Degree	Size	RHS	Direct ms	BSR8	MAGMA best	Best library
128	4	8	8	0.013	1.398	1.299	0.916
128	4	32	8	0.016	0.853	1.769	0.853
128	4	8	64	0.015	0.799	2.297	0.799
128	4	32	64	0.025	1.277	1.863	0.961
128	16	8	8	0.025	0.527	1.001	0.527
128	16	32	8	0.043	0.795	1.221	0.795
128	16	8	64	0.034	0.458	1.874	0.458
128	16	32	64	0.074	1.272	1.719	0.985
128	128	8	8	0.138	0.364	1.031	0.364
128	128	32	8	0.496	0.681	0.724	0.681
128	128	8	64	0.191	0.338	2.407	0.338
128	128	32	64	0.643	1.088	1.783	0.792
512	4	8	8	0.016	0.694	1.848	0.618
512	4	32	8	0.024	1.157	2.254	0.981
512	4	8	64	0.034	0.450	1.705	0.450
512	4	32	64	0.070	1.252	1.825	0.932
512	16	8	8	0.031	0.483	2.204	0.483
512	16	32	8	0.080	1.221	2.213	0.902
512	16	8	64	0.109	0.278	2.080	0.278
512	16	32	64	0.315	1.207	1.677	0.842
512	128	8	8	0.193	0.314	2.743	0.314
512	128	32	8	0.673	1.075	2.130	1.043
512	128	8	64	0.808	0.224	2.285	0.224
512	128	32	64	2.388	1.406	1.783	1.112

Update paths. On the 256-block-row RHS-32 API test, value replacement plus execution takes 0.051 ms in synchronized host time, versus 0.206 ms for borrowing changed structure, validating its metadata, rebuilding the row order, and executing. These are geometric means of three process medians, with twenty samples each. Values alternate between two pre-existing device payloads on fixed degree-four structure. Structural cases alternate between degree-four and degree-eight matrices already on the GPU. Thus this is an API-path cost comparison, not an equal-work kernel speedup or a measurement of assembling a new structure.

3 Historical development evidence

This appendix reports archived development measurements. They are separate from the verified-build campaigns in the main text and are not combined with them as additional repetitions or used to adjust current speedups.

3.1 RHS-8 mapping ablation

The RHS-8 ablation covers 144 configurations at 256, 512, 1,024, and 4,096 block rows and columns. At each size, 32 configurations combine degrees 2, 4, 8, 16, the four shape distributions, and both locality patterns. Four additional configurations combine high-variance or bimodal shapes with both localities and repeating degrees 0, 1, 16, 16.

Each configuration uses matrix seeds 1, 2, and 3, with one fresh process per seed. These are different instances, not repeated processes of one fixed instance. The 432 processes retain 25,920 raw timed calls. Each method receives five warmups and 30 measured calls. Method order is randomized within each process, and jobs are shuffled before serial execution. Input generation, CUDA-event placement, and synchronized host timing use the same definitions as the main campaign. For each configuration, we average the three paired log ratios of GPU medians, then give each configuration equal weight. No significance interval is assigned. The comparator is the scalar RHS-8 kernel, and the tiled method executes the dispatch archived with that campaign, including its small-grid fallback.

Table 9: Scalar baseline time divided by optimized direct time for RHS 8. Each size contains 36 configurations and three matrix seeds. The 256-row paths execute the same scalar kernel, so their deviation from parity reflects measurement variation.

Block rows	RHS-8 mapping	Speedup
256	Original scalar	0.995
512	Eight-row tiles	1.278
1,024	Eight-row tiles	1.491
4,096	Eight-row tiles	2.002

Table 9 shows increasing benefit with grid size. All 108 configurations using the tiled path improve after averaging across seeds, with geometric speedup 1.563 and minimum 1.028. Across all 144 configurations, including the unchanged fallback, the ratio is 1.396. The 4,096-row bimodal/random/degree-16 case improves from 2.245 to 1.086 ms, a ratio of 2.067. These absolute times are geometric means of the three seed medians. The minimum is a descriptive observation, not a statistical guarantee against regression.

Every process passes 64 CPU output probes per method on its timed matrix. A deterministic full-output test exercises all 64 supported block shapes in the tiled path, with empty rows, a non-default stream, and output initialized to NaNs to detect missing writes. That campaign’s updated suite passed memcheck with zero errors and racecheck filtered to the new kernel with zero errors or warnings.

For the representative 4,096-row bimodal/random/degree-16 case, Nsight Compute reports 4,096 CTAs instead of 8,192 and 47 registers per thread instead of 40. Eligible warps per scheduler increase from 0.40 to 0.55, and issued warps per scheduler increase from 0.26 to 0.32. These workload-specific counters support improved scheduling. Instrumented durations do not enter the speedup estimates. The combined experiment does not isolate the contributions of tiling, accumulator reuse, and CTA count. Its paired direct-kernel measurements are separate from the main library comparisons and use three different seeds.

3.2 Earlier width specialization and profiling

The original five-seed dataset remains available under `data/publication/`. It contains 128 configurations per seed, one benchmark process per configuration and seed, five warmups, and 20 measured executions. Its saved summaries reproduce geometric ratios of 4.988 over default CSR and 4.222 over the original grouped composition. That composition rebuilt pointers and cleared all outputs on every call. The original harness also used constant dense input and fixed method order.

Historical timings remain separate from the current comparison. Software, inputs, timing protocol, and comparators changed, so differences cannot isolate one optimization. Earlier standard errors describe five whole-grid averages, not process noise. The configuration bootstrap measures sensitivity to sampling the chosen grid, not uncertainty over an application population. The kernel policy was developed using separate single-seed sweeps with 1,024 block rows, three warmups, and seven measured executions. Table 10 reports those exploratory comparisons. Each ratio compares the named successive variants on the stated slice, with equal configuration weights.

Table 10: Exploratory development ablations, separate from the main evaluation.

Change	Comparator	Cases	Ratio
Wider accumulation	Four-accumulator direct	128	1.24
Cooperative staging	Wider accumulation	128	1.18
Wider accumulation and staging	Four-accumulator direct	128	1.46
RHS-32 buffering and dispatch	Single-buffer staging	32	1.11

Historical Nsight Compute measurements found 99.54% occupancy but only 0.48 eligible warps per scheduler on an early RHS-64 case. Staging reduced long-scoreboard share from 78.0% to 37.1%, while barrier waiting reached 36.4%. The tall RHS-32 profile confirmed native asynchronous copies. These workload-specific observations support the mechanism. Instrumented durations are excluded from speedup calculations.